


Musikla: Language for Generating Musical Events

Pedro M. Silva 

Dummy University Computing Laboratory, Portugal

My second affiliation, Country

<http://www.myhomepage.edu>

johnqpublic@dummyuni.org

José João Almeida 

Algoritmi, Departamento de Informática, Universidade do Minho, Braga, Portugal

jj@di.uminho.pt

Abstract

In this paper, we'll discuss a simple approach to integrating musical events, such as notes or chords, into a programming language. First we'll analyze the problem and its particular requirements. Then we will discuss the solution we developed to meet those requirements. Finally we'll analyze the result and discuss possible alternative routes we could've taken.

2012 ACM Subject Classification Computing methodologies → Language resources

Keywords and phrases Umbundu, Angola Languages, Morphological Analysis, Spell Checking

Digital Object Identifier 10.4230/OASICS.SLATE.2020.23

Funding *Pedro M. Silva*: (Optional) author-specific funding acknowledgements

1 Introduction

Musikla stands for Music and Keyboard Language. Our goal is to develop a DSL (*Domain Specific Language*) that allows treating musical events as regular data in a programming language. More than generating these musical events offline, we want to be able to easily declare keyboards that map keys and expressions to mutate the state or play musical events (or both).

The project can be sectioned in three different, modular layers: inputs, the language, and outputs. While music events can be described as code literals in our language, they can also originate from many other sources. After being processed by our language, it then emits a stream of musical events to the **Player**, which then multiplexes those events into however many outputs the user defined.

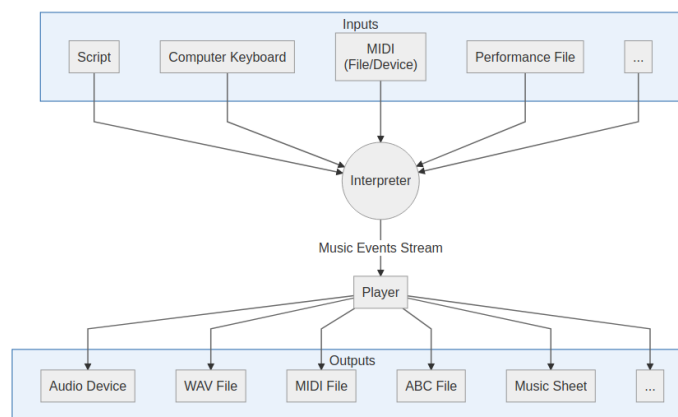


Figure 1 The three main layers of the project.



© Pedro Miguel Silva, José João Almeida;
licensed under Creative Commons License CC-BY

42nd Conference on Very Important Topics (SLATE 2020).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:10

OpenAccess Series in Informatics



OASICS Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

While the development of both the input and output layers, as well as their many respective components presents by itself many interesting challenges that could be discussed, we will instead focus this paper on the aspects of the middle layer: the *interpreter*, while acknowledging the existence (and their effects) of the layers that wrap around it.

As such, the problem of developing the interpreter can be divided into two parts: the syntax used for describing the notes and the operators that compose them inside the language; and the semantics of the generated events, how they are stored in memory, and how their temporal properties are handled without forcing the programmer/user to manually type them.

Designing the syntax for describing those musical expressions, especially given our strong desire to make those musical expressions first class citizens, like other primitive data types in most programming languages (such as numbers, strings or arrays are), did unearth some challenges. To minimize the learning curve for new users, and avoid reinventing the wheel, we decided to adopt a subset of the very popular note declaration syntax from the ABC Notation project[5] and integrate it with our language.

As for the execution model, we decided to go with a tree-walker interpreter[2]. Although computationally slower than other alternatives (such as a bytecode virtual machine), the ease of implementation allowed us to prototype and develop features extremely fast. And with a more mature and stable language in the future, there is always the potential to rewrite the interpreter if performance or latency ever reveal themselves as potential problems.

The simplest way of generating such musical events in a programming language is to use already common, *low-level*, programming mechanisms, such as using a procedural approach where the user creates each event manually by calling a function and providing as parameters all the events' information, such as its timestamp and duration. This is the approach used by some of the existing languages in this space, such as *SonicPi*.

■ **Listing 1** Example of a hypothetical imperative API for creating events

```
play_note( 0, 100, 'A' );
play_note( 100, 50, 'B' );
play_note( 150, 200, 'C' );
```

Instead we decided to follow a more *functional* approach, with custom syntax and operators, as well as the ability to describe those events in a single expression. Musical events are treated as sequences, and as such can be stored in variables, passed around inside functions and transformed. So, for musical events, we will be exploring a way to define them in code, as *musical literals*.

■ **Listing 2** Our proposed declarative syntax that calculates timings implicitly

```
play( A B/2 C2 );
```

2 The Problem and its Requirements

There are two important requirements we need to consider when evaluating possible solutions to this problem: the ability to produce music interactively, and to produce music lazily.

The first requirement, **Interactivity**, relates to our goal of not only being able to generate music offline, but also in a live environment: give the user the ability to program several snippets of musical events, and then control them through a virtual keyboard or through other interactive means.

The second requirement, **Laziness**, refers to a concept that is familiar in functional programming languages: values are generated when we need them, not earlier. In our case,

this implies that a musical sequence could be potentially infinite (like an infinite repetition of some arrangement). If playing this music live, the musician could determine to stop this arrangement sooner or later.

Given these two requirements, we can conclude we **cannot** generate all music events at the start and then sort them to play them in order. Because of that, the events must always be sorted already.

► **Lemma 1** (Total Order). *All operators must return a sequence of events that respects our time unit's total order.*

Data Model

The basic premise is that expressions can generate a special data type: **Music**. Music is simply a sequence (or stream) of ordered musical events.

A musical **Event** can be one of many things, such as a *note*, a *chord*, or even more implementation-specific events like MIDI messages. While all events must have a start time, some events can be instantaneous (events with a duration of zero time units).

The time unit used does not need to be a common time measure, like seconds or milliseconds, and can be really anything so long as it has a **total order**.

Operators

Operators are special operations defined at the syntactic level that allow *music* to be composed in different ways, such as concatenated, parallelized or repeated. Many of these operators can have equivalent functions available through the language that provide more customization (such as a parallel function that stops when the smallest operand stops, instead of the longest).

```

Concatenation Music1 Music2 ... MusicN
  type List[Music] -> Music
Parallel Music1 | Music2 | ... | MusicN
  type List[Music] -> Music
Repetition Music * Integer
  type Music, Integer -> Music
Arpeggio Chord * Music
  type Chord, Music -> Music
Transpose Music + Integer and Music - Integer
  type Music, Integer -> Music

```

It is also useful to establish that while most operators work on sequences of musical events, they can also accept a singular event as their argument: one event can be trivially converted into a sequence of one element. Such occurrence is so common and trivial that the conversion should therefore be implicit whenever necessary.

Grids

Another type available in our language are grids. Also known in most music applications as the process of quantization [4]. The reason it is so useful in our language is that when receiving input as musical events from a live keyboard, their timings are naturally more prone to having small discrepancies that can become more apparent when we then mix them with generated musical events (which have precise timings).

23:4 Musikla: Language for Generating Musical Events

119 Having events always aligned with such a grid can also make computations and trans-
120 formations of such events easier and simpler, which is always a plus for our language.

```
121 Create a grid Grid(size)  
122     type Fraction -> Grid  
123 Aligning grid::align(music)  
124     type Grid, Music -> Music  
125 Compose Grids Grid::compose(grid1, grid2, ..., gridN)  
126     type List[Grid] -> Grid
```

127 We can see that apart from the basic operations of creating a grid and aligning events to
128 said grid, we also want the ability to compose multiple grids (of different precisions). We
129 will approach this matter in more detail later.

130 Keyboards

131 A core part of the language is our hability to declare keyboards, which we can describe as
132 mappings between **Keys** and *Musical Expressions*.

133 Each expression can mutate the state (changing variables or calling functions), return
134 some music (sequence of musical events) to be played, or both.

135 Some of the operations we want to be able to perform with keyboards are as follows:

```
136 Create a keyboard keyboards\create()  
137     type () -> Keyboard  
138 Binding a Key keyboard::register(key, expression)  
139     type Keyboard, Key, Expression -> Keyboard  
140 Mapping a keyboard keyboard::map(transformer)  
141     type Keyboard, ( Music -> Music ) -> Keyboard  
142 Aligning with a Grid keyboard::with_grid(grid)  
143     type Keyboard, Grid -> Keyboard
```

144 3 Implementation

145 The reference implementation for this system is written in Python, although the approach
146 here should be language agnostic.

147 One of the features that Python boasts (but are certainl not exclusive to it) that have
148 eased our implementation are generators[3]. They integrate very nicely into both our concept
149 of emitting musical events as sequences (or iterators, as they are called in Python), as well
150 as into our concept of laziness, where events are generated on demand when needed, and
151 thus infinite musical sequences can be handled easily.

152 Context State

153 To keep track of the *cursor* (the current timestamp where the next event should start) each
154 operator in our language is implemented as a function call that receives an implicit **Context**
155 object. While here we'll mostly focus just on the methods related to time management
156 provided by the context, it can be used to store other types of information, like the default
157 length of a musical note, for instance, to avoid forcing the user to type it out all the time.

158 It is important to keep in mind that there might be more than one context in execution
159 at the same time. This can be most obvious with the use of the parallel operator, where
160 each operand must run concurrently (and thus could not share the same context).

161 Let's describe what kinds of functionality our context should provide.

162 **cursor(ctx)** Return the current cursor position

163 **seek(ctx, time)** Advance the cursor to the given position

164 **fork(ctx)** Clone the parent context and return the new one. Allows multiple concurrent

165 contexts to be used

166 **join(parent, child)** If the child's cursor is ahead, make the parent context catch up

167 3.1 Operators

168 Basic Events

169 The basic building block of our system are the **Note**, **Chord** and **Rest** events. We can

170 use the current *context* to determine the event's timestamp, as well as its default duration

171 (in case the user does not explicitly state one). Any event(s) that is/are not captured in a

172 variable or passed to a function are implicitly played.

173 ■ **Listing 3** Creating a Note Event

```
173 c ' 1/4
```

176 Concatenation

177 We've seen how single events' creation is handled. Now it is important for us to see how

178 we can combine those events together. And probably the most straightforward operator

179 of all, concatenation, it simply consumes each event. Each event, as we've seen before, is

180 responsible for seeking the context depending on the event's duration.

181 ■ **Listing 4** Snippet of Wet Hands by C418

```
181 S4/4 T74 L/8 V90 ;
182 A, E A B ^c B A E D ^F ^c e ^c A3 ;
```



183 ■ **Figure 2** Generated music sheet for concatenation¹.

185 Repetition

186 The repetition operand is in a way very similar to the concatenation operator. It makes sense,

187 since repeating any kind of music pattern *N* times could be thought as a particular case of

188 as concatenation where there are *N* operands, all representing the same musical pattern.

189 ■ **Listing 5** Intro to Westworld's Theme by Ramin Djawadi

```
189 I1 S6/8 T140 L/8 V90 ;
190 A*11 G F*12
```

¹ Rendered with \$ABC_UI. Some hand made changes made for clarity.

23:6 Musikla: Language for Generating Musical Events



■ **Figure 3** Generated music sheet for repetition.

193 Parallel

194 The parallel operator enables playing multiple sequences of musical events simultaneously.
195 However our events are emitted as a single sequence of ordered events, thus requiring merging
196 the multiple sequences into a single one, while maintaining the properties of laziness and
197 order. The operator assumes that each of its operands already maintains those properties
198 on their own, and so is only in charge of making sure the merged sequence does so as well.
199 With this in mind, it relies on a custom *merge sorted* algorithm for iterables (not related to
200 the most common merge sort algorithm by John von Neumann).

■ **Listing 6** Snippet of Soft to Be Strong by Marina

```
201 T120 V70 L1;  
202 r/4 ^g/4 ^g/4 ^g/4 ^f/2 e/8 ^d3/8 ^c2 | [^Cm] [BM] [AM] [BM]  
203
```



■ **Figure 4** Generated music sheet for parallel

205 The merge sorted function receives N operands and creates a buffer with the size N .
206 For each operand it *forks* the context, so that they can execute concurrently and each will
207 mutate their own context only. It then requests one single event for each operand.

208 After the buffer is prefilled (meaning it has at least one event for all non-empty operands),
209 the algorithm finds the earliest event stored in the it. Let's assume it is stored in the K
210 index of the buffer, with $K < N$. The method emits the value stored in `buffer[K]` and then
211 fills requests the next event from the K operand (storing `null` if the operand has no more
212 events to emit). It then repeats this step until all operands have been drained.

213 3.2 Integration in a Programming Environment

214 Apart from generating musical events from somewhat static instructions, our goal is to have
215 those events integrate into a programming language in the same way integers, floats, strings

and booleans do: as data that can be stored, passed around and manipulated. This, of course, must still retain all the properties we've laid out for our sequences of events: being lazy and always being ordered.

Variables and Functions

All expressions that are assigned to a variable run in a forked context, with its cursor set to zero initially. Musical expressions inside variables are still lazy (meaning they only calculate each musical event when the variable is first used, not declared) but the events are cached to prevent the calculations from being performed every time. This cache is then garbage collected when the variable is no longer in use.

Since events declared inside variables have their start time set relative to zero, it needs to be calculated each time the variable is used to replace it with the correct value. This highlights an important aspect of our language: the need for each musical event to be immutable, and any changes made to them to actually be implemented as new instances of the event.

This works well enough because those events are very lightweight objects, and the benefits of not having their values mysteriously changed midway during execution outweigh the small cost of a possible unnecessary allocation of an event that would only be used in one place instead of many.

Function calls, on the other hand, pass the current context to the inside of the function, so that any events played there now their correct times.

When integrating functions into our language, we decided to keep the semantics simple. Emitting musical events inside a function is similar to its return value being an iterator that gives out the emitted events on demand. This means that a value cannot both emit musical events, while also returning other values manually through a return statement.

There is no syntactic marker to distinguish regular functions from musical-emitting ones. Instead, the language runtime starts executing each function as a regular one, and automatically switches its execution mode into a generator-like implementation once the first event is emitted. Any return statements that are evaluated after this point must have no value (thus preserving the feature to early-stop a function). If they do try to return a custom value, a runtime exception is triggered.

Here we can see a small snippet of the beginning of *Fugue 2 in C minor in Book I of the J.S. Bach's Well-Tempered Clavier*, and how using functions and variables can help us visualize the structure behind music.

■ Listing 7 Example of repeating the same note

```

249 fun fugue ( $subj, $resp ) =>
250   ( $subj $resp | stretch( r, $subj ) ( $subj + 7 ) );
251
252
253 S8/4 T140 L/4 V120;
254
255 $subj = r c/2 B/2 c G _A c/2 B/2 c d
256         G c/2 B/2 c d F/2 G/2 _A2 G/2 F/2;
257
258 $resp = E/2 c/2 B/2 A/2    G/2 F/2 _E/2 D/2    C _e d c
259         _B A _B c    ^F G _A F;
260
261 play( fugue( $subj, $resp ) );
262

```



■ **Figure 5** Generated music sheet for fugue example

263 Grids


264 To define a grid there is only one parameter required: the length of it's cells. When aligning
265 musical events, anything inside each cell will be pushed to the closest edge of the cell.

266 Grids are highly customizable too, however. They have multiple parameters, such as
267 **forgiveness** and **range**, that determine when an event is affected by the grid (depending on
268 how close it's start time is to the edge of the cell). Each parameter can even be customized
269 separately for the left and right sides of the cell's edge.

270 Let's take a look at an example of a grid. In this example the grid has a cell size of **1**.
271 We define the same values for both left and right sides just for the sake of this demonstration,
272 but each side could have different values.

■ **Listing 8** Declaring a grid

```
273 $grid = Grid( 1,  
274     forgiveness_left = 125, # 1/8  
275     range_left = 375, # 3/8  
276     forgiveness_right = 125, #1/8  
277     range_right = 375 # 3/8  
278 );  
279  
280
```



■ **Figure 6** Representation of two cells from this grid.

281 We can see in this timeline two cells (each with a size of 1). Any events that fall in the
282 yellow and grey areas are ignored (meaning their timestamps are not changed) while events
283 in the green areas are pushed to whatever edge is closest. But even this behavior can be
284 customized, forcing events to always go to the previous edge cell, or always to the next.

It is then trivial to see how we could compose multiple grids in sequence, each with different ranges (green areas) that capture different events and align them accordingly.

Keyboards

Finally we can combine all the systems we've described above, from musical expressions, grids, variables and functions, and devise a compact way of describing virtual keyboards.

To make the process of designing keyboards less verbose, we've added syntactic sugar to this process, that is translated in the background to regular function calls registering each key binding.

While a picture maybe worth a thousand words, a good example is worth maybe even more. So here we can take a brief look at the workflow for defining two keyboards (that are active at the same time). The first keyboard has all the musical keys (the chords and single notes we want), all aligned by a custom grid.

The second keyboard binds to the up and down arrow keys and allow us to change the virtual instrument through which we play the sounds of the notes in the keyboard (those instruments can be identified by an integer and usually follow the General MIDI standard[1]).

Listing 9 Creating a keyboard that can play multiple instruments

```
$inst = 0;

fun spin_instrument ( ref $instrument, $change ) {
    $instrument = $instrument + $change;

    setinstrument( $instrument );
};

@keyboard {
    a: [^Cm];    s: [BM];    d: [AM];    f: [EM];    g: [^Fm];
    1: ^c;       2: ^d;       3: e;        4: ^f;       5: ^g;
    6: b;        7: ^c';      8: ^d';      9: e';
}::with_grid( Grid( 1 / 16 ) );

@keyboard {
    up: spin_instrument( $inst, 1 );
    down: spin_instrument( $inst, -1 );
};
```

Keyboards are objects (that we could save in a variable for example) and that can perform many operations, like unions and intersections, or maps and filters. They can be enabled and disabled at runtime, and their keys can be simulated to be pressed and released.

More than that, we don't need to restrict ourselves to computer keyboards. We can for instance, define bindings between MIDI events and musical expressions, so that when we connect a piano keyboard to our computer, we can use each piano key to play more than a single note.

Since like we've seen keyboard keys are not limited to computer keyboards, we can imagine the possibilities of events we could listen to: knobs, mouse buttons, the mouse scroll wheel. We could even create an event that could, for example, listen on a socket and trigger when a message is received, allowing in that way our musical applications to be controlled remotely.

The result is that our keyboards are extremely extensible and allow for a great deal of creativity. And thanks to our tight integration with the Python language, those extensions can

334 be easily integrated and don't require hacking the source code or recompiling the application.

335 **4 Results Discussion**

336 To solve our problem of keeping track of the timing implicitly for each event created, we
 337 decided to pass around a context variable. There were other possible solutions, like keeping
 338 this data in some sort of global variable. Our approach does give us some advantages, such
 339 as being able to have multiple contexts in play at the same time. It does have drawbacks,
 340 too. Every function defined in our language must receive this context to be able to create
 341 events at the appropriate time. However, functions defined in Python do not expect this
 342 parameter. Therefore, special conversions must be made when exchanging values between
 343 both languages.

344 Also, our solution to have variables just offset the timings of each event they contain every
 345 time those variables are used simplifies the process of integrating variables into our existing
 346 semantics of music generation. This solution, however, does not answer other questions
 347 unrelated to the timing, such as: should events stored in variables use the musical instrument
 348 set when they were declared, or when the variable was used?

349 However, this early work already provides a solid foundation for a musical *DSL* that
 350 while dynamic (with variables, functions and control structures) integrates very well with
 351 established musical standards such as the MIDI protocol and others.

352 **References**

- 353 **1** Gm 1 sound set. URL: <https://www.midi.org/specifications-old/item/gm-level-1-sound-set>.
- 354 **2** Bob Nystrom. *Crafting Interpreters*. 2020.
- 355 **3** Pep 255 – simple generators, May 2001. URL: <https://www.python.org/dev/peps/pep-0255/>.
- 356 **4** Quantization music. URL: [https://en.wikipedia.org/wiki/Quantization_\(music\)](https://en.wikipedia.org/wiki/Quantization_(music)).
- 357 **5** The abc music standard - the tune body, December 2011. URL: http://abcnotation.com/wiki/abc:standard:v2.1#the_tune_body.
- 358
- 359
- 360